


Against general-purpose frameworks — 256 MiB, single-threaded

PyTorch 2.13.0 / ExecuTorch 1.4.1, each in its native layout. Every shape verified to compute the same function before timing.



A fused implementation of the expensive norm beats a decomposed implementation of the cheap one — the clearest argument in this report for writing the kernel at all.