


Adding threads: our kernel vs the best framework at each count

RMSNorm V7, P-core biased. Solid = 4096x8192 (DRAM-bound); dashed = 128x64, where threading collapses our kernel.



Group granularity caps usable threads at M/256, so our occupancy tops out at 6.6 against PyTorch's 7.8 -- which is why the lead erodes. 96/96 correctness checks passed across all four thread counts.